

1.

$x^t[A]x < 0$  the matrix not semidefinite

$$x^t[A]x = \sum_{i=1}^d \sum_{j=1}^d A^{(i,j)} x^{(i)} x^{(j)} =$$

$$x^t \begin{bmatrix} 1 & 2 \\ 2 & 1 \end{bmatrix} x = \begin{matrix} x^{(1)} \\ x^{(2)} \end{matrix} \begin{bmatrix} 1 & 2 \\ 2 & 1 \end{bmatrix} \begin{matrix} x^{(1)} \\ x^{(2)} \end{matrix} = (x^{(1)})^2 + 4x^{(1)}x^{(2)} + (x^{(2)})^2$$

If  $x^{(1)} = 0$  and  $x^{(2)} = 2$  then  $x^t[A]x$  would be  $< 0$  so it would not be positive semidefinite

$$\text{Consider A to be } \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix} \text{ then } x^t[A]x = \begin{matrix} x^{(1)} \\ x^{(2)} \end{matrix} \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix} \begin{matrix} x^{(1)} \\ x^{(2)} \end{matrix} =$$

$4(x^{(1)})^2 + x^{(1)}x^{(2)} + 4(x^{(2)})^2$  which would be  $> 0$  for all values of  $x^{(1)}$  and  $x^{(2)}$  so our new A would be PSD

2.

$$\text{a. } \frac{d^2}{d\theta_p d\theta_n} \text{ of } n_1 \log(\theta_{y=1}) + (n - n_1) \log(1 - \theta_{y=1}) + \sum_{j=1}^d [r_0^j * \log(\theta_{j|0}) + (n - n_1 - r_0^j) \log(1 - \theta_{j|0}) + r_1^j * \log(\theta_{j|1}) + (n - n_1 - r_1^j) \log(1 - \theta_{j|1})] =$$

$$\begin{aligned} & \sum_{j=1}^d [r_1^j / \theta_{j|1} - (n_1 - r_1^j) / (1 - \theta_{j|1})] d\theta_{j|1} \\ & \sum_{j=1}^d [r_0^j / \theta_{j|0} - (n - n_1 - r_0^j) / (1 - \theta_{j|0})] d\theta_{j|0} \\ & n_1 / (\theta_{y=1}) - (n - n_1) / (1 - \theta_{y=1}) d\theta_{y=1} \end{aligned}$$

A taking the 2d derivative  $\frac{d^2}{d\theta_p d\theta_q}$  where  $p \neq q$  we can see that taking the derivative of any of the above with any theta value not already used in the simplification of the problem would result in 0

$$\text{b. } -f(\theta) = g(\theta) = -n_1 \log(\theta_{y=1}) - (n - n_1) \log(1 - \theta_{y=1}) - \sum_{j=1}^d [r_0^j * \log(\theta_{j|0}) + (n - n_1 - r_0^j) \log(1 - \theta_{j|0}) + r_1^j * \log(\theta_{j|1}) + (n - n_1 - r_1^j) \log(1 - \theta_{j|1})]$$

$$\begin{aligned} & \sum_{j=1}^d [r_1^j / \theta_{j|1} + (n_1 - r_1^j) / (1 - \theta_{j|1})] d\theta_{j|1} \\ & [r_1^j / (\theta_{j|1})^2 + (n_1 - r_1^j) / (1 - \theta_{j|1})^2] \frac{d(2)}{d\theta_{j|1}(2)} \leq 0 \Rightarrow g(\theta) \geq 0 \end{aligned}$$

$$\begin{aligned} & \sum_{j=1}^d [r_0^j / \theta_{j|0} - (n - n_1 - r_0^j) / (1 - \theta_{j|0})] d\theta_{j|0} \\ & -r_0^j / (\theta_{j|0})^2 - (n - n_1 - r_0^j) / (1 - \theta_{j|0})^2 \frac{d(2)}{d\theta_{j|0}(2)} \leq 0 \Rightarrow g(\theta) \geq 0 \end{aligned}$$

$$-n_1/(\theta_{y=1}) + (n - n_1)/(1 - \theta_{y=1}) \frac{d\theta_{y=1}}{d\theta_{y=1}(2)} \leq 0 \Rightarrow g(\theta) \geq 0$$

3.

- a.  $\frac{1}{n} \sum_{i=1}^n (y_i - \theta_i^T x_i)^2 = \frac{1}{n} \sum_{i=1}^n (y_i - h(x))^2 = \frac{1}{n} \sum_{i=1}^n (y_i - (\theta_0^T + \theta^T y_x))^2 \rightarrow$   
 $\frac{1}{n} \sum_{i=1}^n 2(y_i - (\theta_0^T + \theta^T x_i))(-x_i) \frac{d}{d\theta(j)}$
- b.  $\frac{d}{d\theta(i)d\theta(j)} \frac{1}{n} \sum_{j=1}^n 2(y_j - (\theta_0^T + \theta^T x_j))(-x_i) \frac{d}{d\theta(j)} = \frac{1}{n} \sum_{j=1}^n 2(x_j)(-x_i) \frac{d}{d\theta(j)} \frac{d}{d\theta(i)} =$   
 $\frac{2}{n} \sum_{j=1}^n (x_k)(x_k)^T = |x||x|^T$  which is greater than or equal to 0 for all x so the function is psd so it is convex

#####

1. data\_mat=np.insert(data\_mat,0,1.,axis =1)

```
from sklearn.feature_extraction.text import CountVectorizer

## Transform to bag of words representation.
vectorizer = CountVectorizer(analyzer = "word", tokenizer = None, preprocessor = None, stop_words = None, max_features = 4500)
data_features = vectorizer.fit_transform(sents_processed)

print ('The original size: ',data_features.shape)

The original size: (3000, 4500)
```

```
## STUDENT: YOUR CODE STARTS HERE
# Task: Append '1' to the beginning of each vector.
# Hint: You can use data_features.toarray() to transform data_features into a numpy array
# The output should be a numpy array named data_mat

data_mat = data_features.toarray()

data_mat=np.insert(data_mat,0,1.,axis =1)

## STUDENT: CODE ENDS
print ('The updated size: ',data_mat.shape)

The updated size: (3000, 4501)
```

2.

## STUDENT: Start of code ###

```
# - yx
# _____
# ((e^yθ^Tx)+1)
```

```
#initialize ld(theta)
derivatives=np.zeros(weights.size)
```

```
for i in range(labels.size):
    #dot product of weights with each column of feature matrix
    column_product=np.dot(weights,feature_matrix[i,:])
```

```

temp=labels[i]*column_product
value = 1/(1+np.exp(temp))
#take the sum for every road
derivatives = derivatives+(- labels[i]*feature_matrix[i,:]*value)

```

```

return derivatives

```

```

# End of code ###

```

```

def weight_derivative(weights, feature_matrix, labels):
    # Input:
    # weights: weight vector w, a numpy vector of dimension d
    # feature_matrix: numpy array of size n by d, where n is the number of data points, and d is the feature dimension
    # labels: true labels y, a numpy vector of dimension d, each with value -1 or +1
    # Output:
    # Derivative of the regression cost function with respect to the weight w, a numpy array of dimension d

    ## STUDENT: Start of code ###

    # - yx
    #  $\frac{-yx}{(e^{y\theta^T x} + 1)}$ 

    #initialize ld(theta)
    derivatives=np.zeros(weights.size)

    for i in range(labels.size):
        #dot product of weights with each column of feature matrix
        column_product=np.dot(weights,feature_matrix[i,:])
        temp=labels[i]*column_product
        value = 1/(1+np.exp(temp))
        #take the sum for every road
        derivatives = derivatives+(- labels[i]*feature_matrix[i,:]*value)

    return derivatives

# End of code ###

```

```

# STUDENT: PRINT THE OUTPUT AND COPY IT TO THE SOLUTION FILE
my_weights = np.ones(data_mat.shape[1]) # a weight of all 1s
derivative = weight_derivative(my_weights,train_data,train_labels)

print (derivative[:10])

[ 1.23415330e+03 -4.13993755e-08  1.00000000e+00  9.99993856e-01
 1.99987630e+00  9.99859072e-01  9.52574127e-01  3.59772270e+01
 2.99996572e+00 -1.38879439e-11]

```

3.

```

temp_weight = [0.5]*(len(train_data[0]))
initial_weights = np.array(temp_weight)
#print(initial_weights)

```

```

step_size = 5
tolerance = 2
# end of code

```

```

#Initialize the weights, step size and tolerance
# Start of code
#STUDENT: Specify the initial_weights, step_size, and tolerance

temp_weight = [0.5]*(len(train_data[0]))
initial_weights = np.array(temp_weight)
#print(initial_weights)

step_size = 1
tolerance = 1
# end of code

# Use the regression_gradient_descent function to calculate
final_weights = gradient_descent(train_data,train_labels, initial_weights, step_size, tolerance)

# end of code
print ("Here are the final weights after convergence:")
print (final_weights)

Iteration: 573 gradient_magnitude: 3.6682392304496507
Iteration: 574 gradient_magnitude: 2.213136375610675
Iteration: 575 gradient_magnitude: 1.9803704159457487
Iteration: 576 gradient_magnitude: 1.8783445398539662
Iteration: 577 gradient_magnitude: 1.8131520656107953
Iteration: 578 gradient_magnitude: 1.7906836176702727
Iteration: 579 gradient_magnitude: 1.7802176657823705
Iteration: 580 gradient_magnitude: 1.7763035017143993
Iteration: 581 gradient_magnitude: 1.774618305284305
Iteration: 582 gradient_magnitude: 1.7739270604126407
Iteration: 583 gradient_magnitude: 1.7735934127788686
Iteration: 584 gradient_magnitude: 1.7733865413467367
Iteration: 585 gradient_magnitude: 1.7731850853899929
Iteration: 586 gradient_magnitude: 1.7728986819359303
Iteration: 587 gradient_magnitude: 1.7723980736553682
Iteration: 588 gradient_magnitude: 1.771442601239768
Iteration: 589 gradient_magnitude: 1.7695302631054801
Iteration: 590 gradient_magnitude: 1.7655867273725743
Iteration: 591 gradient_magnitude: 1.7573107327870083
Iteration: 592 gradient_magnitude: 1.740345037915812
Iteration: 593 gradient_magnitude: 1.710658997369197
Iteration: 594 gradient_magnitude: 1.6765140698262924

Iteration: 595 gradient_magnitude: 1.6515548064717083
Iteration: 596 gradient_magnitude: 1.6338579665542021
Iteration: 597 gradient_magnitude: 1.619519775674744
Iteration: 598 gradient_magnitude: 1.605520519064422
Iteration: 599 gradient_magnitude: 1.5814491956242431
Iteration: 600 gradient_magnitude: 1.4789406846914526
Iteration: 601 gradient_magnitude: 1.114876333894972
Iteration: 602 gradient_magnitude: 0.8802789949735388
Here are the final weights after convergence:
[ 9.36895823e-02  1.38697761e+01 -2.95000490e+01 ... -1.67588608e+02
 5.00000000e-01 -3.14998337e+01]

```

4.

## STUDENT: CODE STARTS HERE

## Pull out the parameters (theta\_0, theta) of the logistic regression model

```
theta = final_weights #gradient_descent(train_data,train_labels, initial_weights, step_size,
tolerance)
```

```
theta0 = theta[0]
```

```
theta = np.delete(theta,0)
```

## STUDENT: CODE ENDS HERE

5.

```
## STUDENT: YOUR CODE HERE
```

```
predict_array = []
```

```
for i in range(len(feature_matrix)):
```

```
    dot = np.dot(feature_matrix[i],weights)
```

```
    z = -1*dot
```

```
    predict = 1/(1+pow(np.e,z))
```

```
    predict_array.append(predict)
```

```
return np.array(predict_array)
```

```
## STUDENT: CODE ENDS
```

```
def model_predict(feature_matrix,weights):  
    # Prediction made by Logistic regression  
  
    # Input:  
    # feature_matrix: numpy array of size n by d+1  
    #                  note we have included the du  
    # weights: weight vector to start with, a numpy  
    # Output:  
    # Labels: predicted labels, a numpy vector of  
  
    ## STUDENT: YOUR CODE HERE  
    predict_array = []  
  
    for i in range(len(feature_matrix)):  
        dot = np.dot(feature_matrix[i],weights)  
  
        z = -1*dot  
        predict = 1/(1+pow(np.e,z))  
        predict_array.append(predict)  
    return np.array(predict_array)  
    ## STUDENT: CODE ENDS
```

#not test error did not work as written so I wrote my own

```

# STUDENT: copy the output of this section to the solution file
def getError(preds_train,train_labels):
    errs_train = 0
    for i in range(len(preds_train)):
        if((preds_train[i] > 0.0) and (train_labels[i] < 0.0)):
            errs_train+=1
        if((preds_train[i] < 0.0) and (train_labels[i] > 0.0)):
            errs_train+=1
    return errs_train

## Get predictions on training and test data
preds_train = model_predict(train_data,final_weights)
#print(preds_train)
preds_test = model_predict(test_data,final_weights)

## Compute errors

errs_train = np.sum((preds_train > 0.0) is not (train_labels > 0.0))
errs_test = np.sum((preds_test > 0.0) is not (test_labels > 0.0))

error_train = getError(preds_train,train_labels)
error_test = getError(preds_test,test_labels)

print ("Training error: ", float(errs_train)/len(train_labels))
print ("Test error: ", float(errs_test)/len(test_labels))

print ("Training error: ", float(error_train)/len(train_labels))
print ("Test error: ", float(error_test)/len(test_labels))

Training error:  0.0004
Test error:  0.002
Training error:  0.1028
Test error:  0.184

```

6.

```

## STUDENT: YOUR CODE HERE
model_predict_arr = model_predict(feature_matrix, weights)
count = 0;
for i in range(len(model_predict_arr)):
    if (model_predict_arr[i]>0) and (model_predict_arr[i]<1):
        if(model_predict_arr[i]<0.5-gamma) or (model_predict_arr[i]>0.5+gamma):
            count+=1

return count

```

```
def margin_counts(feature_matrix, weights, gamma):
    ## Return number of points for which  $Pr(y=1)$  lies in  $[0, 0.5 - \text{gamma})$  or  $(0.5 + \text{gamma}, 1]$ 

    # Input:
    # feature_matrix: numpy array of size n by d+1, where n is the number of data points,
    #                 note we have included the dummy feature as the first column of the
    # weights: weight vector to start with, a numpy vector of dimension d+1
    # gamma: the margin value
    # Output:
    # number of points for which  $Pr(y=1)$  lies in  $[0, 0.5 - \text{gamma})$  or  $(0.5 + \text{gamma}, 1]$ 


    ## STUDENT: YOUR CODE HERE
    model_predict_arr = model_predict(feature_matrix, weights)
    count = 0;
    for i in range(len(model_predict_arr)):
        if (model_predict_arr[i]>0) and (model_predict_arr[i]<1):
            if (model_predict_arr[i]<0.5-gamma) or (model_predict_arr[i]>0.5+gamma):
                count+=1

    return count

    ## STUDENT: CODE ENDS

    gammas = np.arange(0,0.5,0.01)
    f = np.vectorize(lambda g: margin_counts(test_data, final_weights,g))
    plt.plot(gammas, f(gammas)/500.0, linewidth=2, color='green')
    plt.xlabel('Margin', fontsize=14)
    plt.ylabel('Fraction of points above margin', fontsize=14)
    plt.show()
```

C:\Users\kuent\Anaconda3\lib\site-packages\ipykernel\_launcher.py:19: RuntimeWarning: divide by zero encountered in divide



7.

## STUDENT: YOUR CODE HERE

```
model_predict_arr = model_predict(feature_matrix, weights)
denominator=len(feature_matrix)
numerator = 1
```

```
for i in range(len(model_predict_arr)):
    if (model_predict_arr[i]>0) and (model_predict_arr[i]<1):
        if (model_predict_arr[i]<(0.5-gamma)) or (model_predict_arr[i]>(0.5+gamma)):
            denominator+=1
        if (labels[i]<(0.5-gamma)) and (model_predict_arr[i]<(0.5-gamma)):
            numerator+=1
        if (labels[i]>(0.5+gamma)) and (model_predict_arr[i]>(0.5+gamma)):
```

```
numerator+=1
```


```
return numerator/denominator
```

```
## STUDENT: YOUR CODE ENDS
```

```
def margin_errors(feature_matrix, labels, weights, gamma):  
    ## Return error of predictions that lie in intervals [0, 0.5 - gamma) and (0.5 + gamma, 1]  
  
    # Input:  
    # feature_matrix: numpy array of size n by d+1, where n is the number of data points, d is the number of features  
    # note we have included the dummy feature as the first column of the feature matrix  
    # Labels: true labels y, a numpy vector of dimension n  
    # weights: weight vector to start with, a numpy vector of dimension d+1  
    # gamma: the margin value  
    # Output:  
    # error of predictions that lie in intervals [0, 0.5 - gamma) and (0.5 + gamma, 1]  
  
    ## STUDENT: YOUR CODE HERE  
    model_predict_arr = model_predict(feature_matrix, weights)  
    denominator=len(feature_matrix)  
    numerator = 1  
  
    for i in range(len(model_predict_arr)):  
        if (model_predict_arr[i]>0) and (model_predict_arr[i]<1):  
            if(model_predict_arr[i]<(0.5-gamma)) or (model_predict_arr[i]>(0.5+gamma)):  
                denominator+=1  
                if(labels[i]<(0.5-gamma)) and (model_predict_arr[i]<(0.5-gamma)):  
                    numerator+=1  
                if(labels[i]>(0.5+gamma)) and (model_predict_arr[i]>(0.5+gamma)):  
                    numerator+=1  
  
    return numerator/denominator  
  
    ## STUDENT: YOUR CODE ENDS
```

```
## Plot the result  
plt.plot(gammas, f(gammas), linewidth=2)  
plt.ylabel('Error rate', fontsize=14)  
plt.xlabel('Margin', fontsize=14)  
plt.show()
```

C:\Users\kuent\Anaconda3\lib\site-packages\ipykernel\_launcher



8.

```
## STUDENT: YOUR CODE HERE
```



#the first index is the largest value and the last index is the smallest value  
#the value that is going to be excluded first from the array is in the back

```
weights= np.delete(final_weights,0)
print(len(weights))
large_positive = []
small_negative = []

for i in range(len(weights)):
    #initialize array
    if (len(large_positive)<10):
        large_positive.append([vocab[i],weights[i]])
        small_negative.append([vocab[i],weights[i]])
    else:
        #larger than smallest value in array
        if (final_weights[i]>large_positive[0][1]):
            #remove small add large
            large_positive.pop(0)
            large_positive.append([vocab[i],weights[i]])
            large_positive.sort(key=lambda x: x[1])
        #smaller than largest value in array
        if (final_weights[i]<small_negative[9][1]):
            #remove large add small
            small_negative.pop(9)
            small_negative.append([vocab[i],weights[i]])
            small_negative.sort(key=lambda x: x[1])

#remove weights from print statement
final_positive = []
final_negative = []
for i in range(10):
    final_positive.append(large_positive[i][0])
    final_negative.append(small_negative[i][0])

print('Top Ten')
print(large_positive)
print('Bottom Ten')
print(small_negative)

## STUDENT: CODE ENDS
```

```

## STUDENT: YOUR CODE HERE
#the first index is the largest value and the last index is the smallest value
#the value that is going to be excluded first from the array is in the back

weights= np.delete(final_weights,0)
print(len(weights))
large_positive = []
small_negative = []

for i in range(len(weights)):
    #initialize array
    if (len(large_positive)<10):
        large_positive.append([vocab[i],weights[i]])
        small_negative.append([vocab[i],weights[i]])
    else:
        #larger than smallest value in array
        if (final_weights[i]>large_positive[0][1]):
            #remove small add large
            large_positive.pop(0)
            large_positive.append([vocab[i],weights[i]])
            large_positive.sort(key=lambda x: x[1])
        #smaller than largest value in array
        if (final_weights[i]<small_negative[9][1]):
            #remove large add small
            small_negative.pop(9)
            small_negative.append([vocab[i],weights[i]])
            small_negative.sort(key=lambda x: x[1])

#remove weights from print statement
final_positive = []
final_negative = []
for i in range(10):
    final_positive.append(large_positive[i][0])
    final_negative.append(small_negative[i][0])

print('Top Ten')
print(large_positive)
print('Bottom Ten')
print(small_negative)

## STUDENT: CODE ENDS

```

Each is accompanied with their theta value

```

Top Ten
[['you', 522.3124682700212], ['well', 1198.4122353355472], ['family', 1267.4999997273922], ['happy', 1507.5298332164937], ['fri
endly', 1609.6197732734986], ['amazing', 1756.0169505164492], ['works', 1953.3096254243846], ['fantastic', 2054.3429738609047],
['loved', 2349.88442359567], ['great', 2382.2088423393316]]
Bottom Ten
[['worst', -2435.1562271465036], ['disappointment', -2236.3075883660986], ['waste', -2118.928162274578], ['disappointing', -182
0.0104102997275], ['avoid', -1655.739111789458], ['bland', -1609.8454208770834], ['fails', -1431.4384892762353], ['aren', -132
2.6989221304673], ['difficult', -1264.999999999803], ['writing', -824.9671458584996]]

```